

House Price Prediction using Machine Learning

Prajwal Chittori (2K18/CO/249)

Pranav Roy (2K18/CO/250)

Dept. of Computer Science and Engineering)



Problems faced while buying a house

Problems faced during buying a house:

- 1) Buyers are generally not aware of factors that influence the house prices, Such as-
 - i) No. of storeys.
 - ii) Area outside the house
 - iii) Geographic location and connectivity
- 2) Hence real estate agents are trusted with the communication between buyers and sellers as well as laying down a legal contract for the transfer. This just creates a middle man and increases the cost of houses.



Agenda

- Our goal for this project was to use regression and classification techniques in order to estimate the sale price of a house in Bengaluru, India given the feature and pricing data for around 21,000 houses sold within one year.



Dataset

- ? The Dataset is acquired from Kaggle:
- ? (<https://www.kaggle.com/amitabhajoy/bengaluru-house-price-data>)
- ? There are 10 total attributes in the dataset, four of which we derived from current columns. We are planning to use all 10 attributes in our models.



TOOLS USED




— Machine Learning Classification

Supervised Learning

In supervised learning, we are given a data set and already know what our correct output should look like, having the idea that there is a relationship between the input and the output. Supervised learning problems are categorized into "regression" and "classification" problems.

Unsupervised Learning

Unsupervised learning allows us to approach problems with little or no idea what our results should look like. We can derive structure where we maybe don't know effect of variables by clustering the data based on relationships among the variables in the data. With unsupervised learning there is no feedback based on the prediction results



STEPS



Data Collection

First step is to collect data from different sources and then merge them to form the training dataset.

Source-Kaggle



Training the model

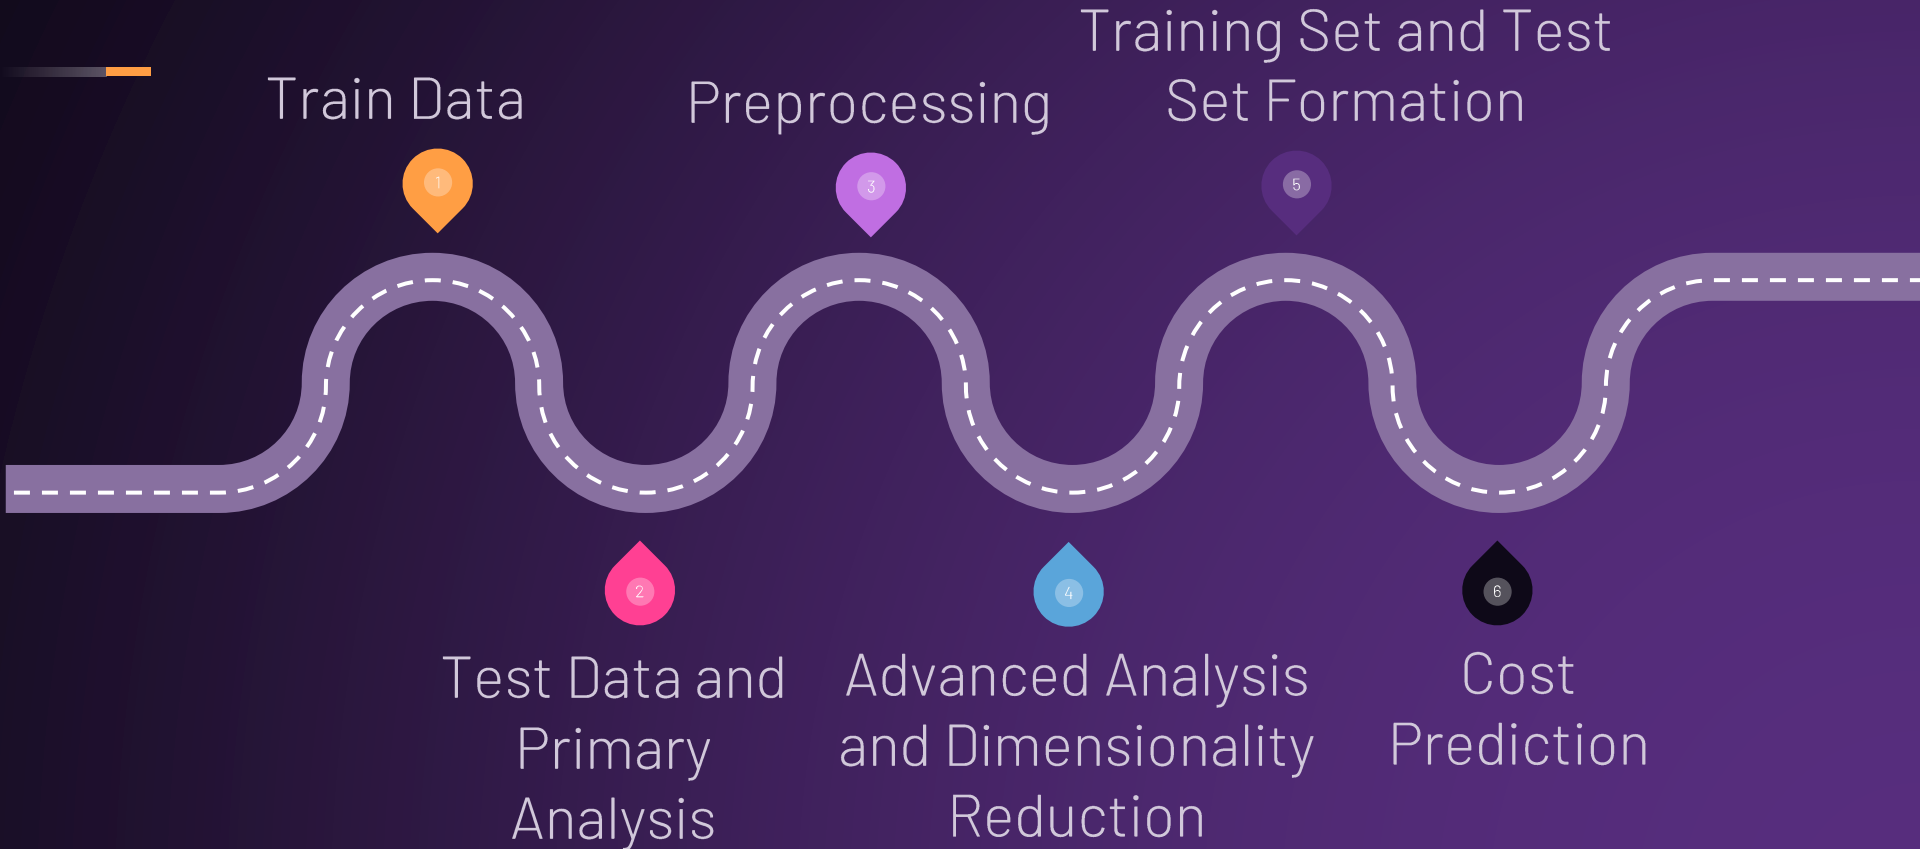
We Use Multiple Linear Regression to train the Dataset. Multiple regression, is a statistical technique that uses several explanatory variables to predict the outcome of a response variable.




Cost Prediction

Based on the generated graphs and trained data we will predict the cost of the house.

Architectural Roadmap



Multiple Linear Regression



- Iteration
 - *Select one variable at a time from the variable importance table created using random forest*
- Significance
 - *Check significance of new variable along with existing variables by its t-value and probability of t- statistics.*
 - *If R-square is improved, keep the variable, else drop it*
- Multicollinearity
 - *Multicollinearity is checked at each step – ensuring the maximum value is < 4*
 - *If new variable has multicollinearity above threshold value – drop it*
 - *If introducing the new variable increases the multicollinearity of any existing variable – then the variable with lowest t-value is dropped*
- Model Accuracy
 - *After adding the new variable, the model accuracy on train and test data using error rate and MAPE is checked.*
 - *Drop the new variable if the model accuracy falls.*





Dataset Details



```
In [19]: housing.info()
```

```
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 20640 entries, 0 to 20639  
Data columns (total 10 columns):  
longitude                20640 non-null float64  
latitude                 20640 non-null float64  
housing_median_age       20640 non-null float64  
total_rooms              20640 non-null float64  
total_bedrooms           20433 non-null float64  
population               20640 non-null float64  
households               20640 non-null float64  
median_income            20640 non-null float64  
median_house_value       20640 non-null float64  
ocean_proximity          20640 non-null object  
dtypes: float64(9), object(1)  
memory usage: 1.6+ MB
```

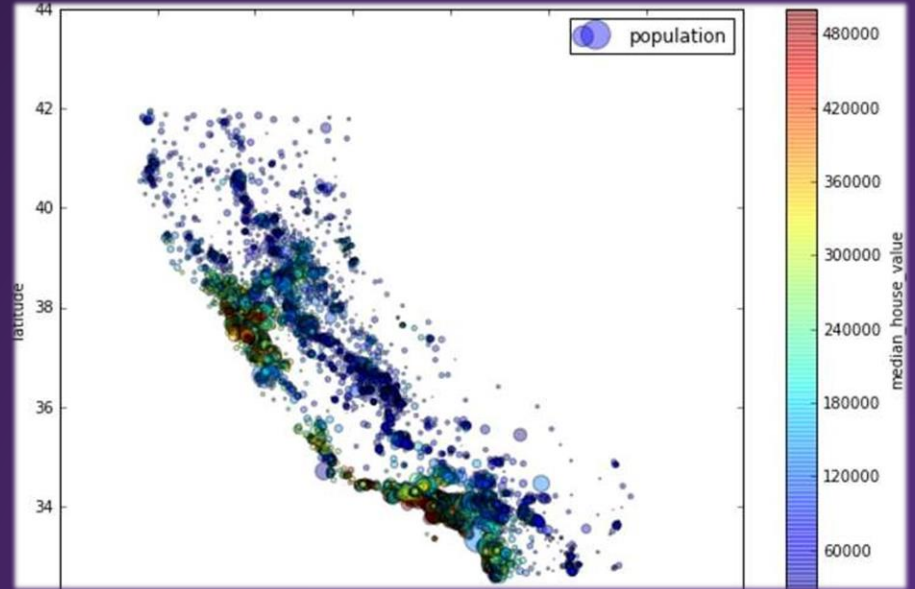
Visualizing using Plot Analysis

Exploratory
Plot
Analysis

Train
Validate
Test

Correlation Plots

Median
Analysis




Plot to visualize role of latitude, longitude & population on the price of the house

Training Dataset Code

```
1 # Multiple Liner Regression
2 from sklearn.linear_model import LinearRegression
3 regressor = LinearRegression()
4 regressor.fit(X_train, y_train)
5 #evaluate the model (intercept and slope)
6 print(regressor.intercept_)
7 print(regressor.coef_)
8 #predicting the test set result
9 y_pred = regressor.predict(X_test)
10 #put results as a DataFrame
11 coeff_df = pd.DataFrame(regressor.coef_, Data.drop
12     ('price',axis =1).columns, columns=['Coefficient'])
13 coeff_df
```

Visualizing Residuals

```
1 # visualizing residuals
2 fig = plt.figure(figsize=(10,5))
3 residuals = (y_test- y_pred)
4 sns.distplot(residuals)
```






Comparing with Actual Results (Code)




```
1 #compare actual output values with predicted values
2 y_pred = regressor.predict(X_test)
3 df = pd.DataFrame({'Actual': y_test, 'Predicted': y_pred})
4 df1 = df.head(10)
5 df1
6 # evaluate the performance of the algorithm (MAE - MSE - RMSE)
7 from sklearn import metrics
8 print('MAE:', metrics.mean_absolute_error(y_test, y_pred))
9 print('MSE:', metrics.mean_squared_error(y_test, y_pred))
10 print('RMSE:', np.sqrt(metrics.mean_squared_error(y_test, y_pred)))
11 print('VarScore:', metrics.explained_variance_score(y_test, y_pred))
```

Mean Absolute Error: 96667.89
Mean Squared Error: 24912134897.75
Root Mean Squared Error: 157835.78
Variance score is: 0.81



Expected Results



— PERFORMANCE ANALYSIS

STRENGTHS

Results are expected to be fairly accurate with accuracy rate of more than 70%

S

WEAKNESSES

Model can improved further with more time given for Training Dataset and using Keras regression

W

Results can be further improved by removing Multi-collinearity

MULTI COLINEARITY

M

Autocorrelation refers to the degree of correlation between the values of the same variables across different observations in the data.

AUTOCORRELATION

A

Keras Regression

Keras is an open-source software library that provides a Python interface for artificial neural networks. Keras acts as an interface for the TensorFlow library.

We can easily fit the regression data with Keras sequential model and predict the test data. Preparing data, Defining the model, Fitting with **Keras Regressor** (accuracy check and visualizing the results) Fitting with the sequential model (accuracy check and visualizing the results) can all be done with Keras



Training Dataset Code

```
test file.cpp x testfile.py
1 # Creating a Neural Network Model
2 from tensorflow.keras.models import Sequential
3 from tensorflow.keras.layers import Dense, Activation
4 from tensorflow.keras.optimizers import Adam
5 model = Sequential()
6 model.add(Dense(19,activation='relu'))
7 model.add(Dense(19,activation='relu'))
8 model.add(Dense(19,activation='relu'))
9 model.add(Dense(19,activation='relu'))
10 model.add(Dense(1))
```

Training

```
test file.cpp x testfile.py
1 model.fit(x=X_train,y=y_train,
2           validation_data=(X_test,y_test),
3           batch_size=128,epochs=400)
4 model.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
=====		
dense (Dense)	multiple	418
dense_1 (Dense)	multiple	380
dense_2 (Dense)	multiple	380
dense_3 (Dense)	multiple	380
dense_4 (Dense)	multiple	20
=====		

Total params: 1,578

Trainable params: 1,578

Non-trainable params: 0



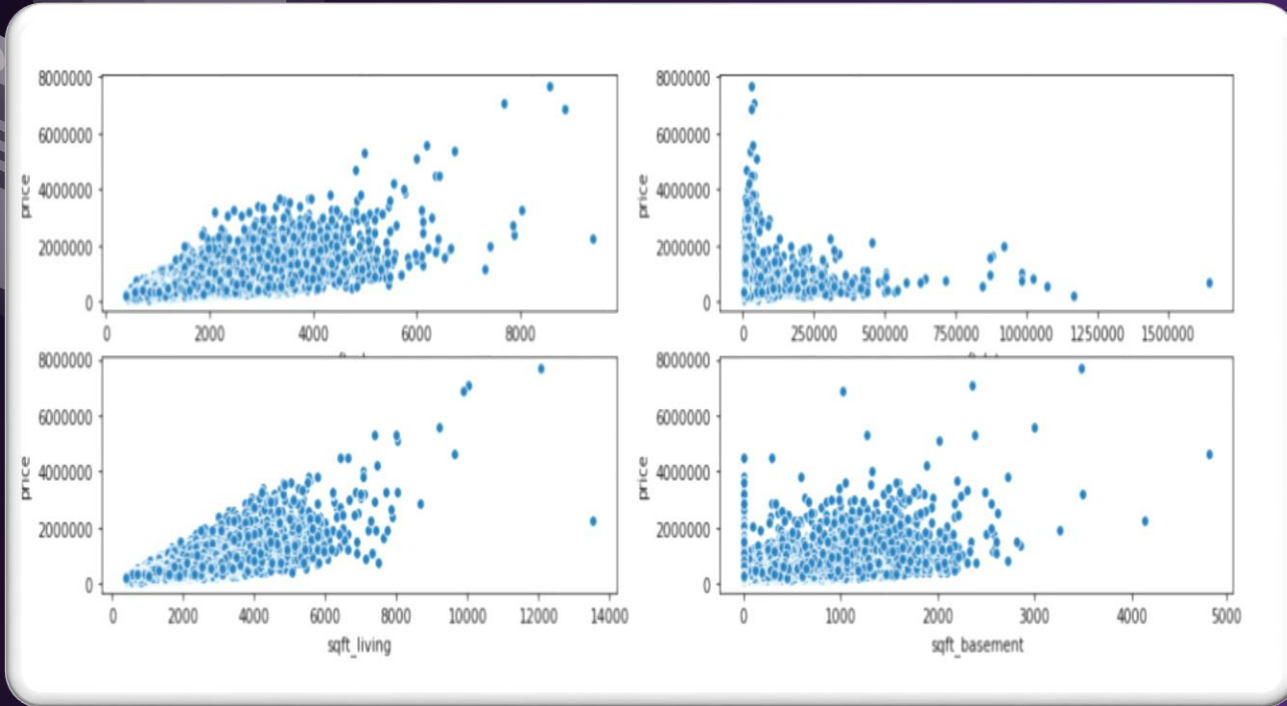
Comparing with Actual Results (Code)



```
test file.cpp  x  testfile.py
1 y_pred = model.predict(X_test)
2 from sklearn import metrics
3 print('MAE:', metrics.mean_absolute_error(y_test, y_pred))
4 print('MSE:', metrics.mean_squared_error(y_test, y_pred))
5 print('RMSE:', np.sqrt(metrics.mean_squared_error(y_test, y_pred)))
6 print('VarScore:', metrics.explained_variance_score(y_test, y_pred))
7 # Visualizing Our predictions
8 fig = plt.figure(figsize=(10,5))
9 plt.scatter(y_test, y_pred)
10 # Perfect predictions
11 plt.plot(y_test, y_test, 'r')
```





Expected Results





Visualising Residuals





— **THANKS!**

